

# DBMS Implementation :

## InnoDB Initialize

woonhak.kang

[woonagi319@skku.edu](mailto:woonagi319@skku.edu)



# Contents

- Before we begin : tools
  - message logging and entry point
  - Coding conventions : the very first
- InnoDB Initialize
  - Init. in-memory objects and check configurations
  - aio system, file, buffer pool, log system
- Create objects
  - lock-system, io-handler, data files and log files
- Open
  - log and system table space
  - transaction system
- Recovery if necessary
- Create double write system if necessary

# Message Log for development

- `ib_logf(log_level, "message")`
  - or you can also do `fprintf(stderr, "")`
- log levels
  - in file : `include/ha_prototypes.h:449`

```
enum ib_log_level_t {  
    IB_LOG_LEVEL_INFO,  
    IB_LOG_LEVEL_WARN,  
    IB_LOG_LEVEL_ERROR,  
    IB_LOG_LEVEL_FATAL  
};
```

# InnoDB entry

- `MYSQL_SRC`
  - e.g.) `/home/woonhak/workspace/mysql-5.6.21`
- `INNODB_SRC`
  - `$MYSQL_SRC/storage/innobase/`
- InnoDB entry point
  - in file: `srv/srv0start.cc:1502`  
`innobase_start_or_create_for_mysql(void)`
  - Starts InnoDB and creates a new database if database files are not found and the user wants.

# Coding conventions

- `srv_XXX` : global system variables
- end of `#if~#else~#endif` directive
- module/function name
  - use abbr.
    - for e.g.)
    - dictionary -> dict, server->srv, utility->ut, memory->mem, buffer->buf
- Assertion
  - `ut_a()`, `ut_ad()`

```
1449 #ifdef UNIV_SYNC_DEBUG
1450 /* Create the thread latch level array where the latch levels
1451    are stored for each OS thread */
1452
1453 sync_thread_level_arrays = static_cast<sync_thread_t*>(
1454     calloc(sizeof(sync_thread_t), OS_THREAD_MAX_N));
1455
1456 ut_a(sync_thread_level_arrays != NULL);
1457
1458 #endif /* UNIV_SYNC_DEBUG */
```

# InnoDB Initialize

- We are under
  - Linux
  - native aio (libaio)
  - use O\_DIRECT
  - does not use raw\_disk (files on file system)
  - default database files are already created
- Initialize
  - aio system
  - file system
  - buffer pool
  - log system

# InnoDB Initialize

- `srv/srv0start.cc:1502`

```
1496 /*****
1497 Starts InnoDB and creates a new database if database files
1498 are not found and the user wants.
1499 @return DB_SUCCESS or error code */
1500 UNIV_INTERN
1501 dberr_t
1502 innobase_start_or_create_for_mysql(void)
1503 /*=====*/
1504 {
1505     ibool    create_new_db;
1506     lsn_t    min_flushed_lsn;
1507     lsn_t    max_flushed_lsn;
1508 #ifdef UNIV_LOG_ARCHIVE
1509     ulint    min_arch_log_no;
1510     ulint    max_arch_log_no;
1511 #endif /* UNIV_LOG_ARCHIVE */
1512     ulint    sum_of_new_sizes;
1513     ulint    sum_of_data_file_sizes;
1514     ulint    tablespace_size_in_header;
1515     dberr_t  err;
1516     unsigned i;
1517     ulint    srv_n_log_files_found = srv_n_log_files;
1518     ulint    io_limit;
1519     mtr_t    mtr;
1520     ib_bh_t* ib_bh;
1521     ulint    n_recovered_trx;
1522     char     logfilename[10000];
1523     char*    logfile0 = NULL;
1524     size_t   dirnamelen;
1525
1526     if (srv_force_recovery > SRV_FORCE_NO_TRX_UNDO) {
1527         srv_read_only_mode = true;
1528     }
1529
1530     if (srv_read_only_mode) {
1531         ib_logf(IB_LOG_LEVEL_INFO, "Started in read only mode");
1532     }
```

# InnoDB Initialize

- `srv_use_native_aio`
  - `include/srv0srv.cc`

```
142 /* If this flag is TRUE, then we will use the native aio of the
143 OS (provided we compiled Innobase with it in), otherwise we will
144 use simulated aio we build below with threads.
145 Currently we support native aio on windows and linux */
146 UNIV_INTERN my_bool srv_use_native_aio = TRUE;
```

- `srv/srv0start.cc:17xx`
  - Check use native AIO

```
1728 #elif defined(LINUX_NATIVE_AIO)
1729
1730 if (srv_use_native_aio) {
1731     ib_logf(IB_LOG_LEVEL_INFO, "Using Linux native AIO");
1732 }
1733 #else
1734 /* Currently native AIO is supported only on windows and linux
1735 and that also when the support is compiled in. In all other
1736 cases, we ignore the setting of innodb_use_native_aio. */
1737 srv_use_native_aio = FALSE;
1738 #endif /* __WIN__ */
```



# InnoDB Initialize

- `srv/srv0start.cc:17XX`
- file flush method : `srv_unix_file_flush_method`
  - `fsync` : use `fsync()` to flush both log and data - default
  - `o_dsync` : use `O_SYNC` for log, `fsync()` for data
  - `o_direct` : use `O_DIRECT` for data, log is buffered, use `fsync()` to flush log and data
  - `o_direct_no_fsync` : use `O_DIRECT` but skip `fsync` for data

```
1740 if (srv_file_flush_method_str == NULL) {
1741     /* These are the default options */
1742
1743     srv_unix_file_flush_method = SRV_UNIX_FSYNC;
1744
1745     srv_win_file_flush_method = SRV_WIN_IO_UNBUFFERED;
1746 #ifndef __WIN__
1747 } else if (0 == ut_strcmp(srv_file_flush_method_str, "fsync")) {
1748     srv_unix_file_flush_method = SRV_UNIX_FSYNC;
1749
1750 } else if (0 == ut_strcmp(srv_file_flush_method_str, "O_DSYNC")) {
1751     srv_unix_file_flush_method = SRV_UNIX_O_DSYNC;
1752
1753 } else if (0 == ut_strcmp(srv_file_flush_method_str, "O_DIRECT")) {
1754     srv_unix_file_flush_method = SRV_UNIX_O_DIRECT;
1755
1756 } else if (0 == ut_strcmp(srv_file_flush_method_str, "O_DIRECT_NO_FSYNC")) {
1757     srv_unix_file_flush_method = SRV_UNIX_O_DIRECT_NO_FSYNC;
```

# buffer pool size threshold

- `srv/srv0start.cc`
  - if buffer pool size less than threshold than buf pool instances = 1
  - `#define BUF_POOL_SIZE_THRESHOLD (1024 * 1024 * 1024)`

```
1792 #define BUF_POOL_SIZE_THRESHOLD (1024 * 1024 * 1024)
1793     srv_max_n_threads = 1    /* io_ibuf_thread */
1794         + 1 /* io_log_thread */
1795         + 1 /* lock_wait_timeout_thread */
1796         + 1 /* srv_error_monitor_thread */
1797         + 1 /* srv_monitor_thread */
1798         + 1 /* srv_master_thread */
1799         + 1 /* srv_purge_coordinator_thread */
1800         + 1 /* buf_dump_thread */
1801         + 1 /* dict_stats_thread */
1802         + 1 /* fts_optimize_thread */
1803         + 1 /* recv_writer_thread */
1804         + 1 /* buf_flush_page_cleaner_thread */
1805         + 1 /* trx_rollback_or_clean_all_recovered */
1806         + 128 /* added as margin, for use of
1807             InnoDB Memcached etc. */
1808         + max_connections
1809         + srv_n_read_io_threads
1810         + srv_n_write_io_threads
1811         + srv_n_purge_threads
1812         /* FTS Parallel Sort */
1813         + fts_sort_pll_degree * FTS_NUM_AUX_INDEX
1814         * max_connections;
1815
1816     if (srv_buf_pool_size < BUF_POOL_SIZE_THRESHOLD) {
1817         /* If buffer pool is less than 1 GB,
1818            use only one buffer pool instance */
1819         srv_buf_pool_instances = 1;
1820     }
```

# server boot

- Boots the InnoDB server
  - invokes `srv_boot()` at `srv0start.cc:1822`
- `srv/srv0srv.cc:1057,srv_boot()`

```
1053 /**
1054 Boots the InnoDB server. */
1055 UNIV_INTERN
1056 void
1057 srv_boot(void)
1058 /*=====*/
1059 {
1060     /* Transform the init parameter values given by MySQL to
1061      use units we use inside InnoDB: */
1062     srv_normalize_init_values();
1063     /* Initialize synchronization primitives, memory management, and thread
1064      local storage */
1065     srv_general_init();
1066     /* Initialize this module */
1067     srv_init();
1068     srv_mon_create();
1069 }
```

init general variables and objects

# server boot

- `srv/srv0srv.cc:1013, srv_general_init()`

```
1008 /**
1009  Initializes the synchronization primitives, memory system, and the thread
1010  local storage. */
1011 UNIV_INTERN
1012 void
1013 srv_general_init(void)
1014 /*=====*/
1015 {
1016     ut_mem_init();
1017     /* Reset the system variables in the recovery module. */
1018     recv_sys_var_init();
1019     os_sync_init();
1020     sync_init();
1021     mem_init(srv_mem_pool_size);
1022     que_init();
1023     row_mysql_init();
1024 }
```

os\_sync\_mutex init

mutex init

additional memory pool (not buffer pool)

# server boot

- `srv/srv0srv.cc:1057, srv_boot()`

```
1053 /**
1054 Boots the InnoDB server. */
1055 UNIV_INTERN
1056 void
1057 srv_boot(void)
1058 /*=====*/
1059 {
1060     /* Transform the init parameter values given by MySQL to
1061     use units we use inside InnoDB: */
1062
1063     srv_normalize_init_values();
1064
1065     /* Initialize synchronization primitives, memory management, and thread
1066     local storage */
1067
1068     srv_general_init();
1069
1070     /* Initialize this module */
1071
1072     srv_init();
1073     srv_mon_create();
1074 }
```

server init.

# server boot

- `srv/srv0srv.cc:909, srv_init()`

```
905 /**
906  * Initializes the server. */
907 UNIV_INTN
908 void
909 srv_init(void)
910 /*=====*/
911 {
912     ulint n_sys_threads = 0;
913     ulint srv_sys_sz = sizeof(*srv_sys);
914
915     #ifndef HAVE_ATOMIC_BUILTINS
916     mutex_create(server_mutex_key, &server_mutex, SYNC_ANY_LATCH);
917     #endif /* !HAVE_ATOMIC_BUILTINS */
918
919     mutex_create(srv_innodb_monitor_mutex_key,
920                 &srv_innodb_monitor_mutex, SYNC_NO_ORDER_CHECK);
921
922     if (!srv_read_only_mode) {
923         /* Number of purge threads + master thread */
924         n_sys_threads = srv_n_purge_threads + 1;
925
926         srv_sys_sz += n_sys_threads * sizeof(*srv_sys->sys_threads);
927     }
928
929     srv_sys = static_cast<srv_sys_t*>(mem_zalloc(srv_sys_sz));
930
931     srv_sys->n_sys_threads = n_sys_threads;
932 }
```

init. system and  
system threads

# of system threads = purge  
thread + master thread

# server boot

- `srv/srv0srv.cc:909, srv_init()`

```
934  if (!srv_read_only_mode) {
935
936      mutex_create(srv_sys_mutex_key, &srv_sys->mutex, SYNC_THREADS);
937
938      mutex_create(srv_sys_tasks_mutex_key,
939                  &srv_sys->tasks_mutex, SYNC_ANY_LATCH);
940
941      srv_sys->sys_threads = (srv_slot_t*) &srv_sys[1];
942
943      for (ulint i = 0; i < srv_sys->n_sys_threads; ++i) {
944          srv_slot_t* slot = &srv_sys->sys_threads[i];
945
946          slot->event = os_event_create();
947          ut_a(slot->event);
948      }
949
950
951      srv_error_event = os_event_create();
952
953      srv_monitor_event = os_event_create();
954
955      srv_buf_dump_event = os_event_create();
956
957      UT_LIST_INIT(srv_sys->tasks);
958  }
```

mutex for srv\_sys

system threads

create event for synchronization  
and signaling

# server boot

- `srv/srv0srv.cc:909, srv_init()`

```
972  /* Create dummy indexes for infimum and supremum */
973
974  dict_ind_init();
975
976  srv_conc_init();
977
978  /* Initialize some INFORMATION SCHEMA internal structures */
979  trx_i_s_cache_init(trx_i_s_cache);
980
981  ut_crc32_init();
982
983  dict_mem_init();
984 }
```

dictionary index



# server boot

- dict/dict0dict.cc:5759, dict\_ind\_init()

```
5755 /**
5756 Inits dict_ind_redundant and dict_ind_compact.
5757 UNIV_INTERN
5758 void
5759 dict_ind_init(void)
5760 /*=====*/
5761 {
5762     dict_table_t*   table;
5763
5764     /* create dummy table and index for REDUNDANT infimum and supremum */
5765     table = dict_mem_table_create("SYS_DUMMY1", DICT_HDR_SPACE, 1, 0, 0);
5766     dict_mem_table_add_col(table, NULL, NULL, DATA_CHAR,
5767         DATA_ENGLISH | DATA_NOT_NULL, 8);
5768
5769     dict_ind_redundant = dict_mem_index_create("SYS_DUMMY1", "SYS_DUMMY1",
5770         DICT_HDR_SPACE, 0, 1);
5771     dict_index_add_col(dict_ind_redundant, table,
5772         dict_table_get_nth_col(table, 0), 0);
5773     dict_ind_redundant->table = table;
5774
5775     /* create dummy table and index for COMPACT infimum and supremum */
5776     table = dict_mem_table_create("SYS_DUMMY2",
5777         DICT_HDR_SPACE, 1,
5778         DICT_TF_COMPACT, 0);
5779     dict_mem_table_add_col(table, NULL, NULL, DATA_CHAR,
5780         DATA_ENGLISH | DATA_NOT_NULL, 8);
5781     dict_ind_compact = dict_mem_index_create("SYS_DUMMY2", "SYS_DUMMY2",
5782         DICT_HDR_SPACE, 0, 1);
5783     dict_index_add_col(dict_ind_compact, table,
5784         dict_table_get_nth_col(table, 0), 0);
5785     dict_ind_compact->table = table;
5786
5787     /* avoid ut_ad(index->cached) in dict_index_get_n_unique_in_tree */
5788     dict_ind_redundant->cached = dict_ind_compact->cached = TRUE;
5789 }
```

dictionary table

by default : we use  
compact type

# server boot

- `srv/srv0srv.cc:909, srv_init()`

```
972  /* Create dummy indexes for infimum and supremum records */
973
974  dict_ind_init();
975
976  srv_conc_init();
977
978  /* Initialize some INFORMATION SCHEMA internal structures */
979  trx_i_s_cache_init(trx_i_s_cache);
980
981  ut_crc32_init();
982
983  dict_mem_init();
984 }
```

checksum utility



# server boot

- ut/ut0crc32.cc:276, ut\_crc32\_init()

```
271 /**
272  Initializes the data structures used by ut_crc32(). Does not do any
273  allocations, would not hurt if called twice, but would be pointless. */
274 UNIV_INTERN
275 void
276 ut_crc32_init()
277 /*=====*/
278 {
279 #if defined(__GNUC__) && defined(__x86_64__)
280     ib_uint32_t vend[3];
281     ib_uint32_t model;
282     ib_uint32_t family;
283     ib_uint32_t stepping;
284     ib_uint32_t features_ecx;
285     ib_uint32_t features_edx;
286
287     ut_cpuid(vend, &model, &family, &stepping,
288             &features_ecx, &features_edx);
312     if (ut_crc32_sse2_enabled) {
313         ut_crc32 = ut_crc32_sse42;
314     } else {
315         ut_crc32_slice8_table_init();
316         ut_crc32 = ut_crc32_slice8;
317     }
318 }
```

for checksum - use  
crc32 sse inst.

# InnoDB Initialize

- `srv/srv0start.cc`
- total io threads

```
1883  /* If user has set the value of innodb_file_io_threads then
1884  we'll emit a message telling the user that this parameter
1885  is now deprecated. */
1886  if (srv_n_file_io_threads != 4) {
1887      ib_logf(IB_LOG_LEVEL_WARN,
1888              "innodb_file_io_threads is deprecated. Please use "
1889              "innodb_read_io_threads and innodb_write_io_threads "
1890              "instead");
1891  }
1892
1893  /* Now overwrite the value on srv_n_file_io_threads */
1894  srv_n_file_io_threads = srv_n_read_io_threads;
1895
1896  if (!srv_read_only_mode) {
1897      /* Add the log and ibuf IO threads. */
1898      srv_n_file_io_threads += 2;
1899      srv_n_file_io_threads += srv_n_write_io_threads;
1900  } else {
1901      ib_logf(IB_LOG_LEVEL_INFO,
1902              "Disabling background IO write threads.");
1903
1904      srv_n_write_io_threads = 0;
1905  }
```

# of io threads = read + write + 2

# InnoDB Initialize

- `srv/srv0start.cc`
- aio system init : `os_aio_init()`

```
1907 ut_a(srv_n_file_io_threads <= SRV_MAX_N_IO_THREADS);
1908
1909 io_limit = 8 * SRV_N_PENDING_IOS_PER_THREAD;
1910
1911 /* On Windows when using native aio the number of aio requests
1912 that a thread can handle at a given time is limited to 32
1913 i.e.: SRV_N_PENDING_IOS_PER_THREAD */
1914 # ifdef __WIN__
1915 if (srv_use_native_aio) {
1916     io_limit = SRV_N_PENDING_IOS_PER_THREAD;
1917 }
1918 # endif /* __WIN__ */
1919
1920 if (!os_aio_init(io_limit,
1921     srv_n_read_io_threads,
1922     srv_n_write_io_threads,
1923     SRV_MAX_N_PENDING_SYNC_IOS)) {
1924
1925     ib_logf(IB_LOG_LEVEL_ERROR,
1926         "Fatal : Cannot initialize AIO sub-system");
1927
1928     return(DB_ERROR);
1929 }
```

`n_pending_ios_per_thread = 32`

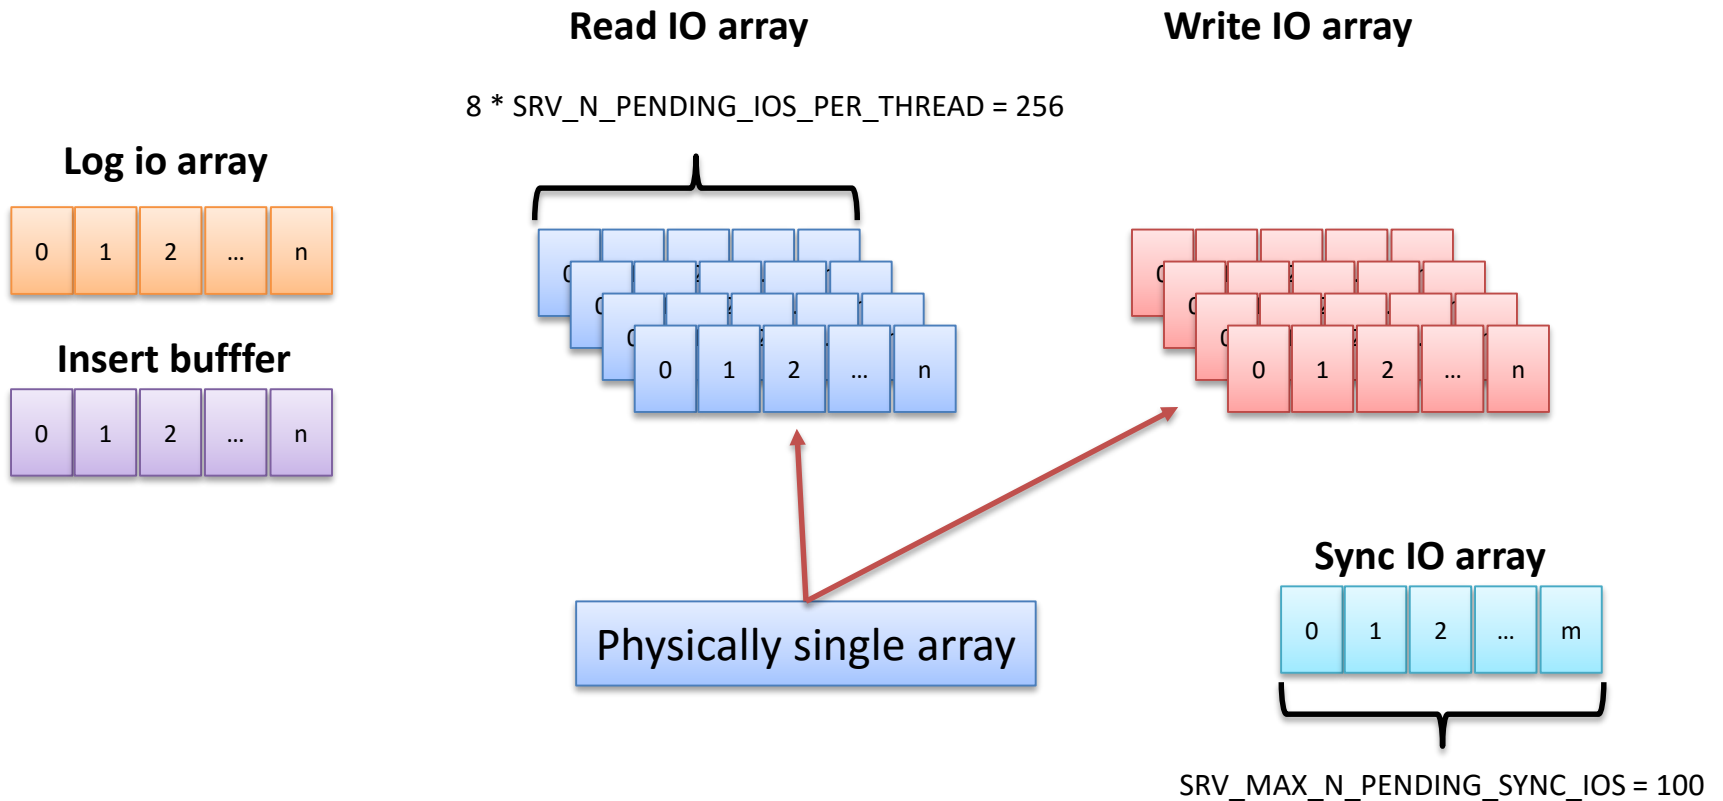
`aio system init`

# aio system init

- Initializes the asynchronous io system.
- Creates one array each for ibuf and log i/o.
- Also creates one array each for read and write where each array is divided logically into `n_read_segs` and `n_write_segs` respectively.
- The caller must create an i/o handler thread for each segment in these arrays.
- This function also creates the sync array.
- No i/o handler thread needs to be created for that

# aio system init

- aio thread and array



# aio system init

- os/os0file.cc:3835, os\_aio\_init()

```
3833 UNIV_INTERN
3834 ibool
3835 os_aio_init(
3836 /*=====*/
3837     uint n_per_seg, /*<! in: maximum number of pending aio
3838         operations allowed per segment */
3839     uint n_read_segs, /*<! in: number of reader threads */
3840     uint n_write_segs, /*<! in: number of writer threads */
3841     uint n_slots_sync) /*<! in: number of slots in the sync aio
3842         array */
3843 {
3844     os_io_init_simple();
3845
3846     #if defined(LINUX_NATIVE_AIO)
3847         /* Check if native aio is supported on this system and tmpfs */
3848         if (srv_use_native_aio && !os_aio_native_aio_supported()) {
3849
3850             ib_logf(IB_LOG_LEVEL_WARN, "Linux Native AIO disabled.");
3851
3852             srv_use_native_aio = FALSE;
3853         }
3854     #endif /* LINUX_NATIVE_AIO */
}
```

check aio \*really\* supported



# aio system init

- os/os0file.cc:3578,  
os\_aio\_native\_aio\_supported()
  - make a temp io context
  - and do test read or write

```
3578 os_aio_native_aio_supported(void)
3579 /*=====*/
3580 {
3581     int      fd;
3582     io_context_t  io_ctx;
3583     char      name[1000];
3584
3585     if (!os_aio_linux_create_io_ctx(1, &io_ctx)) {
3586         /* The platform does not support native aio. */
3587         return(FALSE);
3588     } else if (!srv_read_only_mode) {
3589         /* Now check if tmpdir supports native aio ops. */
3590         fd = innobase_mysql_tmpfile();
3591
3592         if (fd < 0) {
3593             ib_logf(IB_LOG_LEVEL_WARN,
3594                 "Unable to create temp file to check "
3595                 "native AIO support.");
3596
3597             return(FALSE);
3598         }
3599     } else {
```

# aio system init

- os/os0file.cc:3835, os\_aio\_init()
- prepare read io threads

make aio array (kind of io queue)  
see it later

```
3858 os_aio_read_array = os_aio_array_create(  
3859     n_read_segs * n_per_seg, n_read_segs);  
3860  
3861 if (os_aio_read_array == NULL) {  
3862     return(FALSE);  
3863 }  
3864  
3865 uint start = (srv_read_only_mode) ? 0 : 2;  
3866 uint n_segs = n_read_segs + start;  
3867  
3868 /* 0 is the ibuf segment and 1 is the insert buffer segment. */  
3869 for (uint i = start; i < n_segs; ++i) {  
3870     ut_a(i < SRV_MAX_N_IO_THREADS);  
3871     srv_io_thread_function[i] = "read thread";  
3872 }
```

# aio system init

- os/os0file.cc:3835, os\_aio\_init()
- prepare log and ibuf thread

```
3876 if (!srv_read_only_mode) {
3877
3878     os_aio_log_array = os_aio_array_create(n_per_seg, 1);
3879
3880     if (os_aio_log_array == NULL) {
3881         return(FALSE);
3882     }
3883
3884     ++n_segments;
3885
3886     srv_io_thread_function[1] = "log thread";
3887
3888     os_aio_ibuf_array = os_aio_array_create(n_per_seg, 1);
3889
3890     if (os_aio_ibuf_array == NULL) {
3891         return(FALSE);
3892     }
3893
3894     ++n_segments;
3895
3896     srv_io_thread_function[0] = "insert buffer thread";
3897
3898     os_aio_write_array = os_aio_array_create(
3899         n_write_segs * n_per_seg, n_write_segs);
3900
3901     if (os_aio_write_array == NULL) {
3902         return(FALSE);
3903     }
```

# aio system init

- `os/os0file.cc:3835, os_aio_init()`
- prepare write io thread and sync array

```
3898     os_aio_write_array = os_aio_array_create(
3899         n_write_segs * n_per_seg, n_write_segs);
3900
3901     if (os_aio_write_array == NULL) {
3902         return(FALSE);
3903     }
3904
3905     n_segments += n_write_segs;
3906
3907     for (ulint i = start + n_read_segs; i < n_segments; ++i) {
3908         ut_a(i < SRV_MAX_N_IO_THREADS);
3909         srv_io_thread_function[i] = "write thread";
3910     }
3911
3912     ut_ad(n_segments >= 4);
3913 } else {
3914     ut_ad(n_segments > 0);
3915 }
3916
3917 os_aio_sync_array = os_aio_array_create(n_slots_sync, 1);
3918
3919 if (os_aio_sync_array == NULL) {
3920     return(FALSE);
3921 }
```

# aio system init

- os/os0file.cc:3689, create aio array

```
3689 os_aio_array_create(  
3690 /*=====*/  
3691     uint n,      /*!< in: maximum number of pending aio  
3692                  operations allowed; n must be  
3693                  divisible by n_segments */  
3694     uint n_segments) /*!< in: number of segments in the aio array */  
3695 {  
3696     os_aio_array_t* array;  
3697     #ifdef WIN_ASYNC_IO  
3698         OVERLAPPED* over;  
3699     #elif defined(LINUX_NATIVE_AIO)  
3700         struct io_event* io_event = NULL;  
3701     #endif /* WIN_ASYNC_IO */  
3702     ut_a(n > 0);  
3703     ut_a(n_segments > 0);  
3704  
3705     array = static_cast<os_aio_array_t*>(ut_malloc(sizeof(*array)));  
3706     memset(array, 0x0, sizeof(*array));  
3707  
3708     array->mutex = os_mutex_create();  
3709     array->not_full = os_event_create();  
3710     array->is_empty = os_event_create();  
3711  
3712     os_event_set(array->is_empty);  
3713  
3714     array->n_slots = n;  
3715     array->n_segments = n_segments;  
3716  
3717     array->slots = static_cast<os_aio_slot_t*>(  
3718         ut_malloc(n * sizeof(*array->slots)));  
3719  
3720     memset(array->slots, 0x0, sizeof(n * sizeof(*array->slots)));  
3721 }
```

create array struct

create mutex for the array

create slots

# aio system init

- os/os0file.cc:3689, create aio array

```
3725 #if defined(LINUX_NATIVE_AIO)
3726     array->aio_ctx = NULL;
3727     array->aio_events = NULL;
3728
3729     /* If we are not using native aio interface then skip this
3730     part of initialization. */
3731     if (!srv_use_native_aio) {
3732         goto skip_native_aio;
3733     }
3734
3735     /* Initialize the io_context array. One io_context
3736     per segment in the array. */
3737
3738     array->aio_ctx = static_cast<io_context**>(
3739         ut_malloc(n_segments * sizeof(*array->aio_ctx)));
3740
3741     for (ulint i = 0; i < n_segments; ++i) {
3742         if (!os_aio_linux_create_io_ctx(n/n_segments,
3743             &array->aio_ctx[i])) {
3744             /* If something bad happened during aio setup
3745             we should call it a day and return right away.
3746             We don't care about any leaks because a failure
3747             to initialize the io subsystem means that the
3748             server (or atleast the innodb storage engine)
3749             is not going to startup. */
3750             return(NULL);
3751         }
3752     }
3753
3754     /* Initialize the event array. One event per slot. */
3755     io_event = static_cast<struct io_event*>(
3756         ut_malloc(n * sizeof(*io_event)));
3757
3758     memset(io_event, 0x0, sizeof(*io_event) * n);
3759     array->aio_events = io_event;
```

io context

io event array

# aio system init

- `os_aio_array_t`

```
190 /** The asynchronous i/o array structure */
191 struct os_aio_array_t{
192     os_ib_mutex_t mutex; /*!< the mutex protecting the aio array */
193     os_event_t not_full;
194     /*!< The event which is set to the
195         signaled state when there is space in
196         the aio outside the ibuf segment */
197     os_event_t is_empty;
198     /*!< The event which is set to the
199         signaled state when there are no
200         pending i/os in this array */
201     uint n_slots; /*!< Total number of slots in the aio
202         array. This must be divisible by
203         n_threads. */
204     uint n_segments;
205     /*!< Number of segments in the aio
206         array of pending aio requests. A
207         thread can wait separately for any one
208         of the segments. */
209     uint cur_seg; /*!< We reserve IO requests in round
210         robin fashion to different segments.
211         This points to the segment that is to
212         be used to service next IO request. */
213     uint n_reserved;
214     /*!< Number of reserved slots in the
215         aio array outside the ibuf segment */
216     os_aio_slot_t* slots; /*!< Pointer to the slots in the array */
```

```
227 #if defined(LINUX_NATIVE_AIO)
228     io_context_t* aio_ctx;
229     /* completion queue for IO. There is
230         one such queue per segment. Each thread
231         will work on one ctx exclusively. */
232     struct io_event* aio_events;
233     /* The array to collect completed IOs.
234         There is one such event for each
235         possible pending IO. The size of the
236         array is equal to n_slots. */
237 #endif /* LINUX_NATIVE_AIO */
238 };
```

# aio system init

- `os_aio_slot_t`

```
154 /** The asynchronous i/o array slot structure */
155 struct os_aio_slot_t{
156     ibool    is_read; /*!< TRUE if a read operation */
157     ulint    pos;     /*!< index of the slot in the aio
158                        array */
159     ibool    reserved; /*!< TRUE if this slot is reserved */
160     time_t    reservation_time; /*!< time when reserved */
161     ulint    len;     /*!< length of the block to read or
162                        write */
163     byte*    buf;     /*!< buffer used in i/o */
164     ulint    type;     /*!< OS_FILE_READ or OS_FILE_WRITE */
165     os_offset_t offset; /*!< file offset in bytes */
166     os_file_t file;    /*!< file where to read or write */
167     const char* name; /*!< file name or path */
168     ibool    io_already_done; /*!< used only in simulated aio:
169                                TRUE if the physical i/o already
170                                made and only the slot message
171                                needs to be passed to the caller
172                                of os_aio_simulated_handle */
173     fil_node_t* message1; /*!< message which is given by the */
174     void*    message2; /*!< the requester of an aio operation
175                            and which can be used to identify
176                            which pending aio operation was
177                            completed */
183 #elif defined(LINUX_NATIVE_AIO)
184     struct iocb control; /* Linux control block for aio */
185     int    n_bytes; /* bytes written/read. */
186     int    ret; /* AIO return code */
187 #endif /* WIN_ASYNC_IO */
188 };
```



# InnoDB Initialize

- file system init
- srv0start.cc:1931, fil\_init()

invoke file system init

```
1931 fil_init(srv_file_per_table ? 50000 : 5000, srv_max_n_open_files);
1932
1933 double size;
1934 char unit;
1935
1936 if (srv_buf_pool_size >= 1024 * 1024 * 1024) {
1937     size = ((double) srv_buf_pool_size) / (1024 * 1024 * 1024);
1938     unit = 'G';
1939 } else {
1940     size = ((double) srv_buf_pool_size) / (1024 * 1024);
1941     unit = 'M';
1942 }
1943
1944 /* Print time to initialize the buffer pool */
1945 ib_logf(IB_LOG_LEVEL_INFO,
1946     "Initializing buffer pool, size = %.1f%c", size, unit);
1947
1948 err = buf_pool_init(srv_buf_pool_size, srv_buf_pool_instances);
1949
```

# file system init

- `fil_system_t`

```
261 struct fil_system_t {
262 #ifndef UNIV_HOTBACKUP
263     ib_mutex_t    mutex;    /*!< The mutex protecting the cache */
264 #endif /* !UNIV_HOTBACKUP */
265     hash_table_t* spaces;    /*!< The hash table of spaces in the
266                             system; they are hashed on the space
267                             id */
268     hash_table_t* name_hash; /*!< hash table based on the space
269                             name */
270     UT_LIST_BASE_NODE_T(fil_node_t) LRU;
271     /*!< base node for the LRU list of the
272     most recently used open files with no
273     pending i/o's; if we start an i/o on
274     the file, we first remove it from this
275     list, and return it to the start of
276     the list when the i/o ends;
277     log files and the system tablespace are
278     not put to this list: they are opened
279     after the startup, and kept open until
280     shutdown */
281     UT_LIST_BASE_NODE_T(fil_space_t) unflushed_spaces;
282     /*!< base node for the list of those
283     tablespaces whose files contain
284     unflushed writes; those spaces have
285     at least one file node where
286     modification_counter > flush_counter */
287     ulint    n_open;    /*!< number of files currently open */
288     ulint    max_n_open; /*!< n_open is not allowed to exceed
289                             this */
290     ib_int64_t modification_counter; /*!< when we write to a file we
291                                     increment this by one */

```

```
292     ulint    max_assigned_id; /*!< maximum space id in the existing
293                             tables, or assigned during the time
294                             mysqld has been up; at an InnoDB
295                             startup we scan the data dictionary
296                             and set here the maximum of the
297                             space id's of the tables there */
298     ib_int64_t tablespace_version;
299     /*!< a counter which is incremented for
300     every space object memory creation;
301     every space mem object gets a
302     'timestamp' from this; in DISCARD/
303     IMPORT this is used to check if we
304     should ignore an insert buffer merge
305     request */
306     UT_LIST_BASE_NODE_T(fil_space_t) space_list;
307     /*!< list of all file spaces */
308     ibool    space_id_reuse_warned;
309     /* !< TRUE if fil_space_create()
310     has issued a warning about
311     potential space_id reuse */
312 };

```

# file system init

- fil/fil0fil.cc:1676, fil\_init()

```
1672 /**
1673  * Initializes the tablespace memory cache. */
1674 UNIV_INTERN
1675 void
1676 fil_init(
1677 /*=====*/
1678     ulint hash_size, /*!< in: hash table size */
1679     ulint max_n_open) /*!< in: max number of open files */
1680 {
1681     ut_a(fil_system == NULL);
1682
1683     ut_a(hash_size > 0);
1684     ut_a(max_n_open > 0);
1685
1686     fil_system = static_cast<fil_system_t*>(
1687         mem_zalloc(sizeof(fil_system_t)));
1688
1689     mutex_create(fil_system_mutex_key,
1690                 &fil_system->mutex, SYNC_ANY_LATCH);
1691
1692     fil_system->spaces = hash_create(hash_size);
1693     fil_system->name_hash = hash_create(hash_size);
1694
1695     UT_LIST_INIT(fil_system->LRU);
1696
1697     fil_system->max_n_open = max_n_open;
1698 }
```

Protected by mutex

file system hash

Open file LRU

# InnoDB Initialize

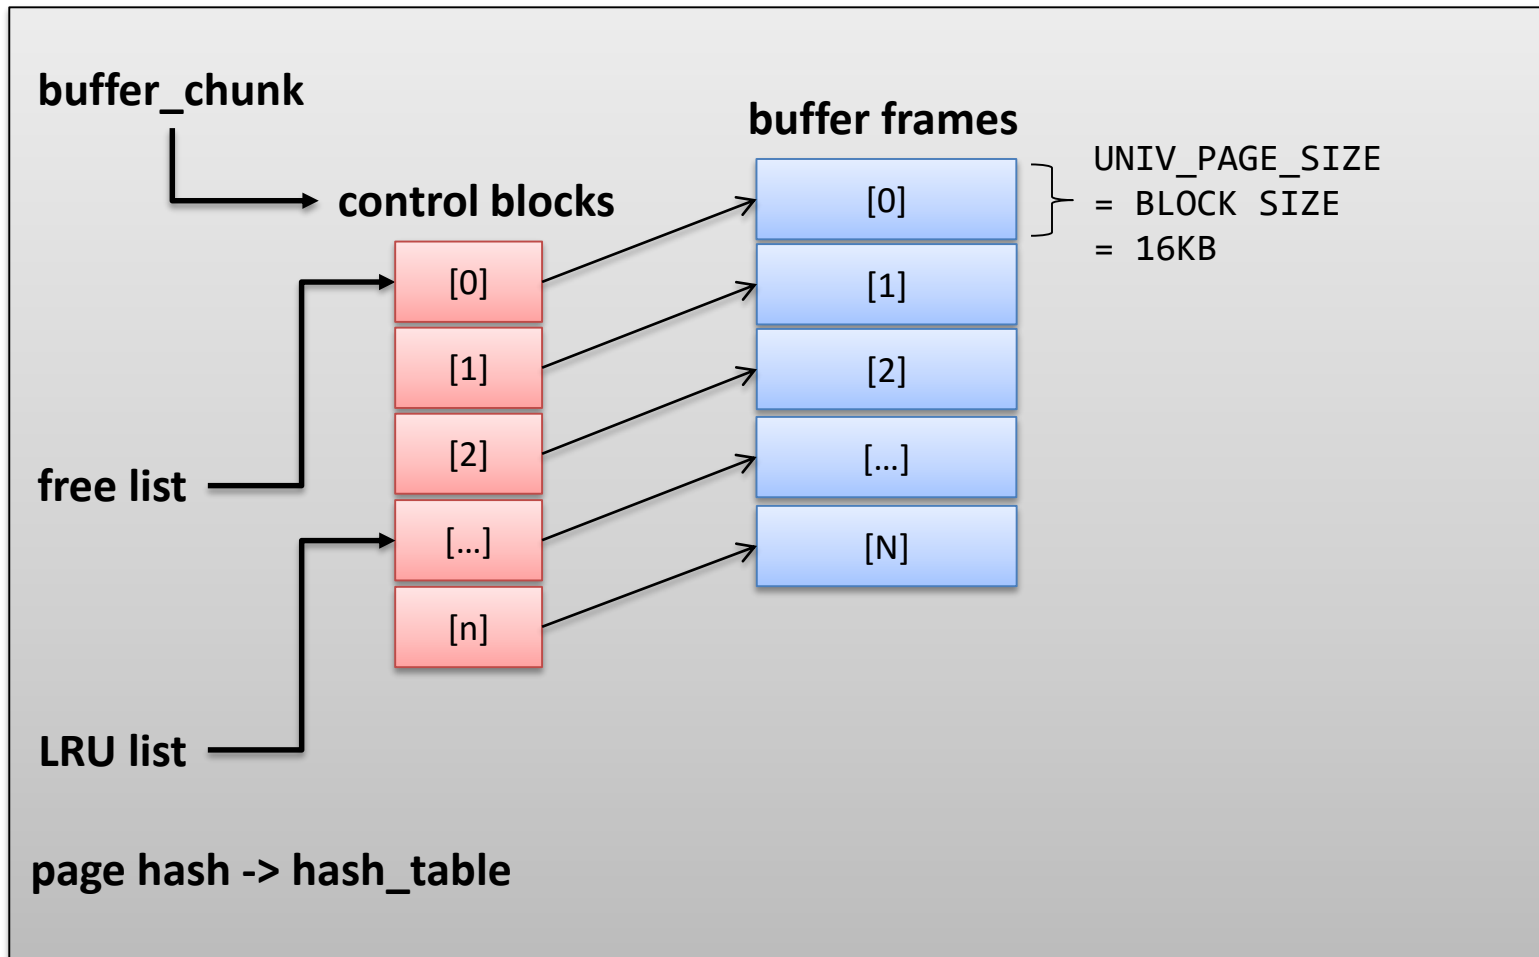
- Buffer Pool Init
- `srv0start.cc:1948, buf_pool_init()`

```
1931  fil_init(srv_file_per_table ? 50000 : 5000, srv_max_n_open_files);
1932
1933  double size;
1934  char unit;
1935
1936  if (srv_buf_pool_size >= 1024 * 1024 * 1024) {
1937      size = ((double) srv_buf_pool_size) / (1024 * 1024 * 1024);
1938      unit = 'G';
1939  } else {
1940      size = ((double) srv_buf_pool_size) / (1024 * 1024);
1941      unit = 'M';
1942  }
1943
1944  /* Print time to initialize the buffer pool */
1945  ib_logf(IB_LOG_LEVEL_INFO,
1946      "Initializing buffer pool, size = %.1f%c", size, unit);
1947
1948  err = buf_pool_init(srv_buf_pool_size, srv_buf_pool_instances);
```

invoke buffer pool init

# Buffer pool init

`buffer_pool_ptr[]`



# Buffer pool init

- buf/buf0buf.cc:1384, buf\_pool\_init()

```
1384 buf_pool_init(  
1385 /*=====*/  
1386   uint total_size, /*!< in: size of the total pool in bytes */  
1387   uint n_instances) /*!< in: number of instances */  
1388 {  
1389   uint i;  
1390   const uint size = total_size / n_instances;  
1391  
1392   ut_ad(n_instances > 0);  
1393   ut_ad(n_instances <= MAX_BUFFER_POOLS);  
1394   ut_ad(n_instances == srv_buf_pool_instances);  
1395  
1396   buf_pool_ptr = (buf_pool_t*) mem_zalloc(  
1397     n_instances * sizeof *buf_pool_ptr);  
1398  
1399   for (i = 0; i < n_instances; i++) {  
1400     buf_pool_t* ptr = &buf_pool_ptr[i];  
1401  
1402     if (buf_pool_init_instance(ptr, size, i) != DB_SUCCESS) {  
1403  
1404       /* Free all the instances created so far. */  
1405       buf_pool_free(i);  
1406  
1407       return(DB_ERROR);  
1408     }  
1409   }  
1410  
1411   buf_pool_set_sizes();  
1412   buf_LRU_old_ratio_update(100 * 3 / 8, FALSE);  
1413  
1414   btr_search_sys_create(buf_pool_get_curr_size() / sizeof(void*) / 64);  
1415  
1416   return(DB_SUCCESS);  
1417 }
```

buffer pool init per instance



# Buffer pool init

- buf/buf0buf.cc:1253, buf\_pool\_init\_instance()

```
1248 /*******  
1249 Initialize a buffer pool instance.  
1250 @return DB_SUCCESS if all goes well. */  
1251 UNIV_INTERN  
1252 ulint  
1253 buf_pool_init_instance(  
1254 /*=====*/  
1255     buf_pool_t* buf_pool, /*!< in: buffer pool instance */  
1256     ulint buf_pool_size, /*!< in: size in bytes */  
1257     ulint instance_no) /*!< in: id of the instance */  
1258 {  
1259     ulint i;  
1260     buf_chunk_t* chunk;  
1261  
1262     /* 1. Initialize general fields  
1263     ----- */  
1264     mutex_create(buf_pool_mutex_key,  
1265                 &buf_pool->mutex, SYNC_BUF_POOL);  
1266     mutex_create(buf_pool_zip_mutex_key,  
1267                 &buf_pool->zip_mutex, SYNC_BUF_BLOCK);  
1268  
1269     buf_pool_mutex_enter(buf_pool);  
1270  
1271     if (buf_pool_size > 0) {  
1272         buf_pool->n_chunks = 1;  
1273  
1274         buf_pool->chunks = chunk =  
1275             (buf_chunk_t*) mem_zalloc(sizeof *chunk);  
1276  
1277         UT_LIST_INIT(buf_pool->free);
```

create mutex and enter

buf\_pool\_zip : use  
compression

# Buffer pool init

- buf/buf0buf.cc:1253,  
buf\_pool\_instance()

```
1277     UT_LIST_INIT(buf_pool->free);
1278
1279     if (!buf_chunk_init(buf_pool, chunk, buf_pool_size)) {
1280         mem_free(chunk);
1281         mem_free(buf_pool);
1282
1283         buf_pool_mutex_exit(buf_pool);
1284
1285         return(DB_ERROR);
1286     }
1287
1288     buf_pool->instance_no = instance_no;
1289     buf_pool->old_pool_size = buf_pool_size;
1290     buf_pool->curr_size = chunk->size;
1291     buf_pool->curr_pool_size = buf_pool->curr_size * UNIV_PAGE_SIZE;
```

init buffer chunk - see this later



# Buffer pool init

- buf/buf0buf.cc:1253,  
buf\_pool\_instance()

```
1293  /* Number of locks protecting page_hash must be a
1294  power of two */
1295  srv_n_page_hash_locks = static_cast<ulong>(
1296  ut_2_power_up(srv_n_page_hash_locks));
1297  ut_a(srv_n_page_hash_locks != 0);
1298  ut_a(srv_n_page_hash_locks <= MAX_PAGE_HASH_LOCKS);
1299
1300  buf_pool->page_hash = ha_create(2 * buf_pool->curr_size,
1301  srv_n_page_hash_locks,
1302  MEM_HEAP_FOR_PAGE_HASH,
1303  SYNC_BUF_PAGE_HASH);
1304
1305  buf_pool->zip_hash = hash_create(2 * buf_pool->curr_size);
1306
1307  buf_pool->last_printout_time = ut_time();
```


create page hash table

create zip page hash table for  
compression

# Buffer pool init

- buffer chunk

```
39 /** A chunk of buffers. The buffer pool is allocated in chunks. */
40 struct buf_chunk_t{
41     uint mem_size; /*!< allocated size of the chunk */
42     uint size; /*!< size of frames[] and blocks[] */
43     void* mem; /*!< pointer to the memory area which
44                was allocated for the frames */
45     buf_block_t* blocks; /*!< array of buffer control blocks */
46 };
47
```

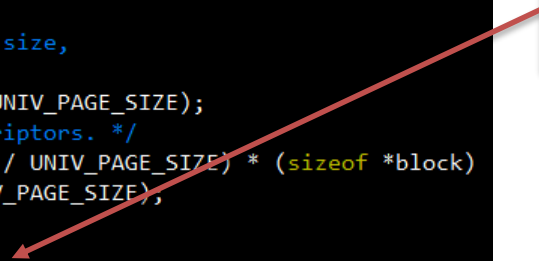


# Buffer pool init

- buf/buf0buf.cc:1041, chunk\_init()

```
1036 /**
1037  * Allocates a chunk of buffer frames.
1038  * @return chunk, or NULL on failure */
1039 static
1040 buf_chunk_t*
1041 buf_chunk_init(
1042     /*-----*/
1043     buf_pool_t* buf_pool, /*!< in: buffer pool instance */
1044     buf_chunk_t* chunk,    /*!< out: chunk of buffers */
1045     ulint mem_size) /*!< in: requested size in bytes */
1046 {
1047     buf_block_t* block;
1048     byte* frame;
1049     ulint i;
1050
1051     /* Round down to a multiple of page size,
1052      * although it already should be. */
1053     mem_size = ut_2pow_round(mem_size, UNIV_PAGE_SIZE);
1054     /* Reserve space for the block descriptors. */
1055     mem_size += ut_2pow_round((mem_size / UNIV_PAGE_SIZE) * (sizeof *block)
1056                             + (UNIV_PAGE_SIZE - 1), UNIV_PAGE_SIZE);
1057
1058     chunk->mem_size = mem_size;
1059     chunk->mem = os_mem_alloc_large(&chunk->mem_size);
1060
1061     if (UNIV_UNLIKELY(chunk->mem == NULL)) {
1062         return(NULL);
1063     }
1064 }
```

allocate chunk mem  
(blocks + frames)



# Buffer pool init

- buf/buf0buf.cc:1041, chunk\_init()

```
1066  /* Allocate the block descriptors from
1067  the start of the memory block. */
1068  chunk->blocks = (buf_block_t*) chunk->mem;
1069
1070  /* Align a pointer to the first frame. Note that when
1071  os_large_page_size is smaller than UNIV_PAGE_SIZE,
1072  we may allocate one fewer block than requested. When
1073  it is bigger, we may allocate more blocks than requested. */
1074
1075  frame = (byte*) ut_align(chunk->mem, UNIV_PAGE_SIZE);
1076  chunk->size = chunk->mem_size / UNIV_PAGE_SIZE
1077    - (frame != chunk->mem);
1078
1079  /* Subtract the space needed for block descriptors. */
1080  {
1081    ulint size = chunk->size;
1082
1083    while (frame < (byte*) (chunk->blocks + size)) {
1084      frame += UNIV_PAGE_SIZE;
1085      size--;
1086    }
1087
1088    chunk->size = size;
1089  }
```

allocate control blocks

allocate frame  
(page size aligned)

# Buffer pool init

- buf/buf0buf.cc:1041, chunk\_init()

```
1091 /* Init block structs and assign frames for them. Then we
1092 assign the frames to the first blocks (we already mapped the
1093 memory above). */
1094
1095 block = chunk->blocks;
1096
1097 for (i = chunk->size; i--; ) {
1098     buf_block_init(buf_pool, block, frame);
1099     UNIV_MEM_INVALID(block->frame, UNIV_PAGE_SIZE);
1100
1101     /* Add the block to the free list */
1102     UT_LIST_ADD_LAST(list, buf_pool->free, (&block->page));
1103
1104     ut_d(block->page.in_free_list = TRUE);
1105     ut_ad(buf_pool_from_block(block) == buf_pool);
1106
1107     block++;
1108     frame += UNIV_PAGE_SIZE;
1109 }
1110
1111
1112 #ifdef PFS_GROUP_BUFFER_SYNC
1113     pfs_register_buffer_block(chunk);
1114 #endif
1115     return(chunk);
```

init control block

add all blocks to free list

# Buffer pool init

- buf/buf0buf.cc:971, buf\_block\_init()

```
967 /**
968  * Initializes a buffer control block when the buf_pool is created. */
969 static
970 void
971 buf_block_init(
972     /*=====*/
973     buf_pool_t* buf_pool, /*!< in: buffer pool instance */
974     buf_block_t* block,   /*!< in: pointer to control block */
975     byte* frame)         /*!< in: pointer to buffer frame */
976 {
977     UNIV_MEM_DESC(frame, UNIV_PAGE_SIZE);
978
979     block->frame = frame;
980
981     block->page.buf_pool_index = buf_pool_index(buf_pool);
982     block->page.state = BUF_BLOCK_NOT_USED;
983     block->page.buf_fix_count = 0;
984     block->page.io_fix = BUF_IO_NONE;
985
986     block->modify_clock = 0;
```

set data frame

page.buf\_pool\_index : back pointer  
page.state : current status  
page.buf\_fix\_count : reference count  
page.io\_fix : block fix  
(Shared/eXclusive)

# Buffer pool init

- buf/buf0buf.cc:1384, buf\_pool\_init()

```
1384 buf_pool_init(  
1385 /*=====*/  
1386  uint_t total_size, /*!< in: size of the total pool in bytes */  
1387  uint_t n_instances) /*!< in: number of instances */  
1388 {  
1389     uint_t i;  
1390     const uint_t size = total_size / n_instances;  
1391  
1392     ut_ad(n_instances > 0);  
1393     ut_ad(n_instances <= MAX_BUFFER_POOLS);  
1394     ut_ad(n_instances == srv_buf_pool_instances);  
1395  
1396     buf_pool_ptr = (buf_pool_t*) mem_zalloc(  
1397         n_instances * sizeof *buf_pool_ptr);  
1398  
1399     for (i = 0; i < n_instances; i++) {  
1400         buf_pool_t* ptr = &buf_pool_ptr[i];  
1401  
1402         if (buf_pool_init_instance(ptr, size, i) != DB_SUCCESS) {  
1403  
1404             /* Free all the instances created so far. */  
1405             buf_pool_free(i);  
1406  
1407             return(DB_ERROR);  
1408         }  
1409     }  
1410  
1411     buf_pool_set_sizes();  
1412     buf_LRU_old_ratio_update(100 * 3 / 8, FALSE);  
1413  
1414     btr_search_sys_create(buf_pool_get_curr_size() / sizeof(void*) / 64);  
1415  
1416     return(DB_SUCCESS);  
1417 }
```

set old ratio : 3/8 -> for mid-point insertion

# log init

- `srv/srv0start:1975, log_init()`

```
1974 fsp_init();
1975 log_init();
1976
1977 lock_sys_create(srv_lock_table_size);
1978
1979 /* Create i/o-handler threads: */
1980
1981 for (i = 0; i < srv_n_file_io_threads; ++i) {
1982     n[i] = i;
1983
1984     os_thread_create(io_handler_thread, n + i, thread_ids + i);
1985 }
1986
1987
1988 #ifdef UNIV_LOG_ARCHIVE
1989 if (0 != ut_strcmp(srv_log_group_home_dir, srv_arch_dir)) {
1990     ut_print_timestamp(stderr);
1991     fprintf(stderr, "InnoDB: Error: you must set the log group home dir in my.cnf\n");
1992     ut_print_timestamp(stderr);
1993     fprintf(stderr, "InnoDB: the same as log arch dir.\n");
1994
1995     return(DB_ERROR);
1996 }
1997 #endif /* UNIV_LOG_ARCHIVE */
1998
```

log init



# log init

- log/log0log.cc:838, log\_init()

```
835 Initializes the log. */
836 UNIV_INTERN
837 void
838 log_init(void)
839 /*=====*/
840 {
841     log_sys = static_cast<log_t*>(mem_alloc(sizeof(log_t)));
842     mutex_create(log_sys_mutex_key, &log_sys->mutex, SYNC_LOG);
843     mutex_create(log_flush_order_mutex_key,
844                 &log_sys->log_flush_order_mutex,
845                 SYNC_LOG_FLUSH_ORDER);
846     mutex_enter(&(log_sys->mutex));
847
848     /* Start the lsn from one log block from zero: this way every
849     log record has a start lsn != zero, a fact which we will use */
850
851     log_sys->lsn = LOG_START_LSN;
852
853     ut_a(LOG_BUFFER_SIZE >= 16 * OS_FILE_LOG_BLOCK_SIZE);
854     ut_a(LOG_BUFFER_SIZE >= 4 * UNIV_PAGE_SIZE);
855
856     log_sys->buf_ptr = static_cast<byte*>(
857         mem_zalloc(LOG_BUFFER_SIZE + OS_FILE_LOG_BLOCK_SIZE));
858
859     log_sys->buf = static_cast<byte*>(
860         ut_align(log_sys->buf_ptr, OS_FILE_LOG_BLOCK_SIZE));
861
862     log_sys->buf_size = LOG_BUFFER_SIZE;
863     log_sys->is_extending = false;
```

create log sys

mutex for protection  
and synchronization

start\_lsn != 0

buf\_ptr : mem buffer  
buf : page aligned buffer for log

# log init

- log/log0log.cc:838, log\_init()

```
879 log_sys->buf_next_to_write = 0;
880
881 log_sys->write_lsn = 0;
882 log_sys->current_flush_lsn = 0;
883 log_sys->flushed_to_disk_lsn = 0;
884
885 log_sys->written_to_some_lsn = log_sys->lsn;
886 log_sys->written_to_all_lsn = log_sys->lsn;
887
888 log_sys->n_pending_writes = 0;
889
890 log_sys->no_flush_event = os_event_create();
891
892 os_event_set(log_sys->no_flush_event);
893
894 log_sys->one_flushed_event = os_event_create();
895
896 os_event_set(log_sys->one_flushed_event);
897
898 /*-----*/
899
900 log_sys->next_checkpoint_no = 0;
901 log_sys->last_checkpoint_lsn = log_sys->lsn;
902 log_sys->n_pending_checkpoint_writes = 0;
903
904
905 rw_lock_create(checkpoint_lock_key, &log_sys->checkpoint_lock,
906                SYNC_NO_ORDER_CHECK);
907
908 log_sys->checkpoint_buf_ptr = static_cast<byte*>(
909     mem_zalloc(2 * OS_FILE_LOG_BLOCK_SIZE));
910
911 log_sys->checkpoint_buf = static_cast<byte*>(
912     ut_align(log_sys->checkpoint_buf_ptr, OS_FILE_LOG_BLOCK_SIZE));
913
```

init log sys variables

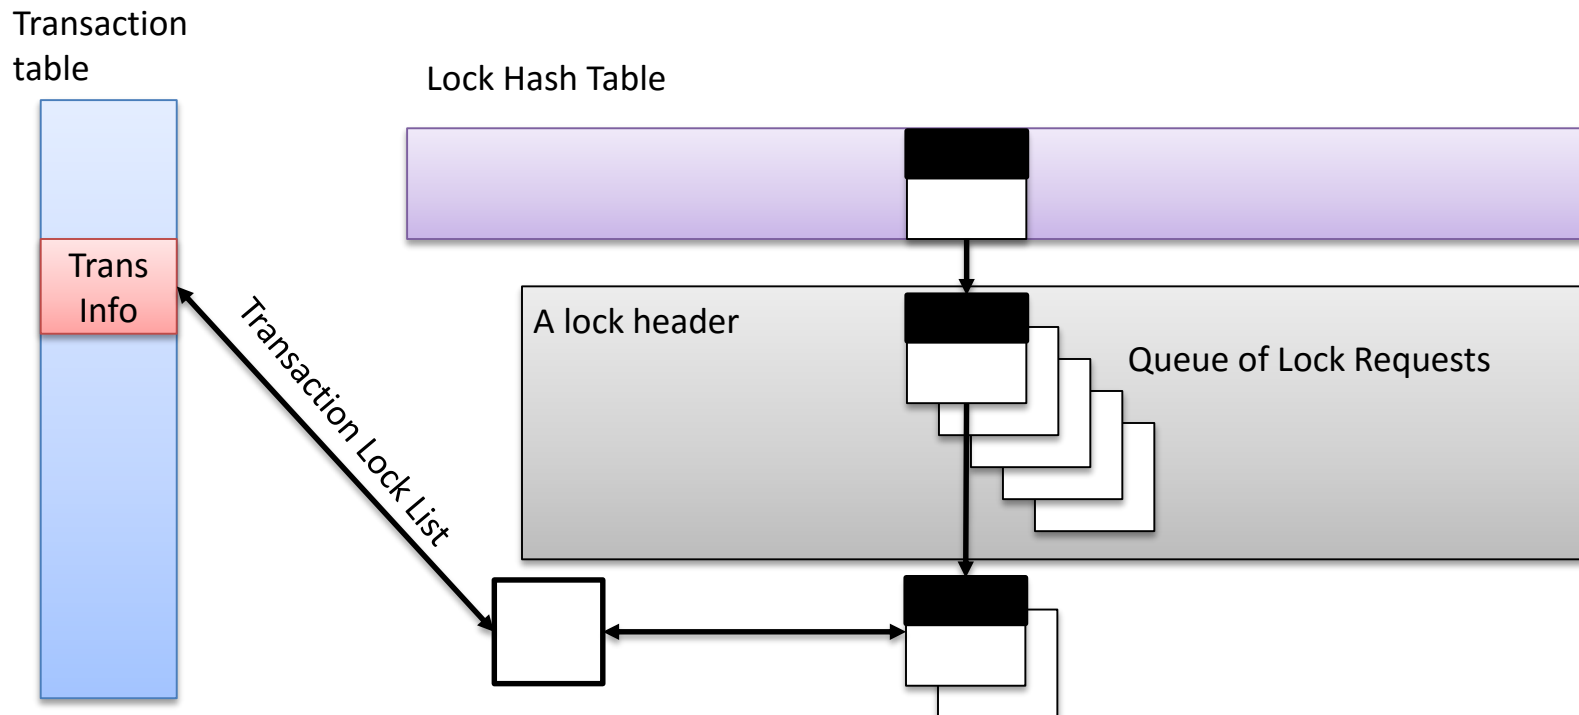
# Contents

- Before we begin : tools
  - message logging and entry point
  - Coding conventions : the very first
- InnoDB Initialize
  - Init. in-memory objects and check configurations
  - aio system, file, buffer pool, log system
- Create objects
  - lock-system, io-handler, data files and log files
- Open
  - log and system table space
  - transaction system
- Recovery if necessary
- Create double write system if necessary

# **CREATE SERVER OBJECTS**

# Locking system

- Locking (and transaction model)
  - When a transaction updates a row in a table, or locks it with SELECT FOR UPDATE, InnoDB establishes a list or queue of locks on that row.
  - Similarly, InnoDB maintains a list of locks on a table for table-level locks transactions hold.



# Locking system

- Lock Compatibility
  - S : shared, X: eXclusive

	Granted Mode						
Requested Mode	None	IS	IX	S	SIX	U	X
IS	+	+	+	+	+	-	-
IX	+	+	+	-	-	-	-
S	+	+	-	+	-	-	-
SIX	+	+	-	-	-	-	-
U	+	-	-	+	-	-	-
X	+	-	-	-	-	-	-

# Locking system

- `srv/srv0start:1977, lock_sys_create()`

```
1974 fsp_init();
1975 log_init();
1976
1977 lock_sys_create(srv_lock_table_size);
1978
1979 /* Create i/o-handler threads: */
1980
1981 for (i = 0; i < srv_n_file_io_threads; ++i) {
1982     n[i] = i;
1983
1984     os_thread_create(io_handler_thread, n + i, thread_ids + i);
1985 }
1986
1987
1988 #ifdef UNIV_LOG_ARCHIVE
1989 if (0 != ut_strcmp(srv_log_group_home_dir, srv_arch_dir)) {
1990     ut_print_timestamp(stderr);
1991     fprintf(stderr, "InnoDB: Error: you must set the log group home dir in my.cnf\n");
1992     ut_print_timestamp(stderr);
1993     fprintf(stderr, "InnoDB: the same as log arch dir.\n");
1994
1995     return(DB_ERROR);
1996 }
1997 #endif /* UNIV_LOG_ARCHIVE */
1998
```

lock system create

# Locking system

- lock table size
  - in `srv_normalize_init_values()`
  - $5 * \text{\# of pages in the buffer pool}$

```
/**
Normalizes init parameter values to use units we use inside InnoDB. */
static
void
srv_normalize_init_values(void)
/*=====*/
{
    ulint n;
    ulint i;

    n = srv_n_data_files;

    for (i = 0; i < n; i++) {
        srv_data_file_sizes[i] = srv_data_file_sizes[i]
            * ((1024 * 1024) / UNIV_PAGE_SIZE);
    }

    srv_last_file_size_max = srv_last_file_size_max
        * ((1024 * 1024) / UNIV_PAGE_SIZE);

    srv_log_file_size = srv_log_file_size / UNIV_PAGE_SIZE;

    srv_log_buffer_size = srv_log_buffer_size / UNIV_PAGE_SIZE;

    srv_lock_table_size = 5 * (srv_buf_pool_size / UNIV_PAGE_SIZE);
}
```

lock table size



# Lock system

- lock/lock0lock:589, lock\_sys\_create()

```
/**
Creates the lock system at database start. */
UNIV_INTERN
void
lock_sys_create(
/*=====*/
  ulint n_cells) /*!< in: number of slots in lock hash table */
{
  ulint lock_sys_sz;

  lock_sys_sz = sizeof(*lock_sys)
    + OS_THREAD_MAX_N * sizeof(srv_slot_t);

  lock_sys = static_cast<lock_sys_t*>(mem_zalloc(lock_sys_sz));

  lock_stack = static_cast<lock_stack_t*>(
    mem_zalloc(sizeof(*lock_stack) * LOCK_STACK_SIZE));

  void* ptr = &lock_sys[1];

  lock_sys->waiting_threads = static_cast<srv_slot_t*>(ptr);

  lock_sys->last_slot = lock_sys->waiting_threads;

  mutex_create(lock_sys_mutex_key, &lock_sys->mutex, SYNC_LOCK_SYS);

  mutex_create(lock_sys_wait_mutex_key,
    &lock_sys->wait_mutex, SYNC_LOCK_WAIT_SYS);

  lock_sys->timeout_event = os_event_create();

  lock_sys->rec_hash = hash_create(n_cells);

  if (!srv_read_only_mode) {
    lock_latest_err_file = os_file_create_tmpfile();
    ut_a(lock_latest_err_file);
  }
}
```

create record hash

# Create io threads

- `srv/srv0start:1985`

```
1974 fsp_init();
1975 log_init();
1976
1977 lock_sys_create(srv_lock_table_size);
1978
1979 /* Create i/o-handler threads: */
1980
1981 for (i = 0; i < srv_n_file_io_threads; ++i) {
1982     n[i] = i;
1983
1984     os_thread_create(io_handler_thread, n + i, thread_ids + i);
1985 }
1986
1987
1988 #ifdef UNIV_LOG_ARCHIVE
1989 if (0 != ut_strcmp(srv_log_group_home_dir, srv_arch_dir)) {
1990     ut_print_timestamp(stderr);
1991     fprintf(stderr, "InnoDB: Error: you must set the log group home dir in my.cnf\n");
1992     ut_print_timestamp(stderr);
1993     fprintf(stderr, "InnoDB: the same as log arch dir.\n");
1994
1995     return(DB_ERROR);
1996 }
1997 #endif /* UNIV_LOG_ARCHIVE */
1998
```

create io threads



# Recovery system

- Recovery system
  - It mainly focuses on write-ahead-log and double write
  - If the system was not clean shutdown, then we have to do recovery

# Recovery system

- `srv/srv0start:2056, recv_sys_create()`

```
2056  recv_sys_create();
2057  recv_sys_init(buf_pool_get_curr_size());
2058
2059  err = open_or_create_data_files(&create_new_db,
2060  #ifdef UNIV_LOG_ARCHIVE
2061      &min_arch_log_no, &max_arch_log_no,
2062  #endif /* UNIV_LOG_ARCHIVE */
2063      &min_flushed_lsn, &max_flushed_lsn,
2064      &sum_of_new_sizes);
2065  if (err == DB_FAIL) {
2066
2067      ib_logf(IB_LOG_LEVEL_ERROR,
2068          "The system tablespace must be writable!");
2069
2070      return(DB_ERROR);
2071  } else if (err != DB_SUCCESS) {
2072
2073      ib_logf(IB_LOG_LEVEL_ERROR,
2074          "Could not open or create the system tablespace. If "
2075          "you tried to add new data files to the system "
2076          "tablespace, and it failed here, you should now "
2077          "edit innodb_data_file_path in my.cnf back to what "
2078          "it was, and remove the new ibdata files InnoDB "
2079          "created in this failed attempt. InnoDB only wrote "
2080          "those files full of zeros, but did not yet use "
2081          "them in any way. But be careful: do not remove "
2082          "old data files which contain your precious data!");
2083
2084      return(err);
2085  }
2086  }
```

create recovery system

# Recovery system

- log/log0recv.cc:196, recv\_sys\_create()
  - just make recv\_sys

```
192 /**
193  * Creates the recovery system. */
194 UNIV_INTERN
195 void
196 recv_sys_create(void)
197 /*=====*/
198 {
199     if (recv_sys != NULL) {
200
201         return;
202     }
203
204     recv_sys = static_cast<recv_sys_t*>(mem_zalloc(sizeof(*recv_sys)));
205
206     mutex_create(recv_sys_mutex_key, &recv_sys->mutex, SYNC_RECV);
207
208     #ifndef UNIV_HOTBACKUP
209         mutex_create(recv_writer_mutex_key, &recv_sys->writer_mutex,
210             SYNC_LEVEL_VARYING);
211     #endif /* !UNIV_HOTBACKUP */
212
213     recv_sys->heap = NULL;
214     recv_sys->addr_hash = NULL;
215 }
```

# Recovery system

- `srv/srv0start:2057, recv_sys_init()`

```
2056  recv_sys_create();
2057  recv_sys_init(buf_pool_get_curr_size());
2058
2059  err = open_or_create_data_files(&create_new_db,
2060 #ifdef UNIV_LOG_ARCHIVE
2061      &min_arch_log_no, &max_arch_log_no,
2062 #endif /* UNIV_LOG_ARCHIVE */
2063      &min_flushed_lsn, &max_flushed_lsn,
2064      &sum_of_new_sizes);
2065  if (err == DB_FAIL) {
2066
2067      ib_logf(IB_LOG_LEVEL_ERROR,
2068          "The system tablespace must be writable!");
2069
2070      return(DB_ERROR);
2071  } else if (err != DB_SUCCESS) {
2072
2073      ib_logf(IB_LOG_LEVEL_ERROR,
2074          "Could not open or create the system tablespace. If "
2075          "you tried to add new data files to the system "
2076          "tablespace, and it failed here, you should now "
2077          "edit innodb_data_file_path in my.cnf back to what "
2078          "it was, and remove the new ibdata files InnoDB "
2079          "created in this failed attempt. InnoDB only wrote "
2080          "those files full of zeros, but did not yet use "
2081          "them in any way. But be careful: do not remove "
2082          "old data files which contain your precious data!");
2083
2084      return(err);
2085  }
2086 }
```

init recovery system

# Recovery system

- Definition recovery system
  - include/log0recv.h:384,recv\_sys\_t

```
383 /** Recovery system data structure */
384 struct recv_sys_t{
385 #ifndef UNIV_HOTBACKUP
386     ib_mutex_t    mutex; /*!< mutex protecting the fields apply_log_recs,
387                          n_addrs, and the state field in each recv_addr
388                          struct */
389     ib_mutex_t    writer_mutex; /*!< mutex coordinating
390                          flushing between recv_writer_thread and
391                          the recovery thread. */
392 #endif /* !UNIV_HOTBACKUP */
393     ibool    apply_log_recs;
394             /*!< this is TRUE when log rec application to
395             pages is allowed; this flag tells the
396             i/o-handler if it should do log record
397             application */
398     ibool    apply_batch_on;
399             /*!< this is TRUE when a log rec application
400             batch is running */
401     lsn_t    lsn; /*!< log sequence number */
402     ulint    last_log_buf_size;
403             /*!< size of the log buffer when the database
404             last time wrote to the log */
405     byte*    last_block;
406             /*!< possible incomplete last recovered log
407             block */
408     byte*    last_block_buf_start;
409             /*!< the nonaligned start address of the
410             preceding buffer */
```

# Recovery system

- Definition recovery system
  - `include/log0recv.h:384,recv_sys_t`

```
411 byte*   buf; /*!< buffer for parsing log records */
412 uint    len; /*!< amount of data in buf */
413 lsn_t    parse_start_lsn;
414          /*!< this is the lsn from which we were able to
415          start parsing log records and adding them to
416          the hash table; zero if a suitable
417          start point not found yet */
418 lsn_t    scanned_lsn;
419          /*!< the log data has been scanned up to this
420          lsn */
421 uint     scanned_checkpoint_no;
422          /*!< the log data has been scanned up to this
423          checkpoint number (lowest 4 bytes) */
424 uint     recovered_offset;
425          /*!< start offset of non-parsed log records in
426          buf */
427 lsn_t     recovered_lsn;
428          /*!< the log records have been parsed up to
429          this lsn */
430 lsn_t     limit_lsn; /*!< recovery should be made at most
431          up to this lsn */
432 ibool     found_corrupt_log;
433          /*!< this is set to TRUE if we during log
434          scan find a corrupt log block, or a corrupt
435          log record, or there is a log parsing
436          buffer overflow */
```



# Recovery system

- Definition recovery system
  - `include/log0recv.h:384,recv_sys_t`

```
437 #ifdef UNIV_LOG_ARCHIVE
438     log_group_t*  archive_group;
439         /*!< in archive recovery: the log group whose
440            archive is read */
441 #endif /* !UNIV_LOG_ARCHIVE */
442     mem_heap_t* heap; /*!< memory heap of log records and file
443        addresses*/
444     hash_table_t* addr_hash; /*!< hash table of file addresses of pages */
445     ulint      n_addrs; /*!< number of not processed hashed file
446        addresses in the hash table */
447
448     recv_dblwr_t  dblwr;
449 };
```

# Recovery system

- log/log0recv.cc:380, recv\_sys\_init()

```
376 /*****
377 Inits the recovery system for a recovery operation. */
378 UNIV_INTERN
379 void
380 recv_sys_init(
381 /*=====*/
382  ulint available_memory) /*!< in: available memory in bytes */
383 {
384   if (recv_sys->heap != NULL) {
385     return;
386   }
387 }
388
389 #ifndef UNIV_HOTBACKUP
390 /* Initialize red-black tree for fast insertions into the
391 flush_list during recovery process.
392 As this initialization is done while holding the buffer pool
393 mutex we perform it before acquiring recv_sys->mutex */
394 buf_flush_init_flush_rbt();
395
396 mutex_enter(&(recv_sys->mutex));
397
398 recv_sys->heap = mem_heap_create_typed(256,
399   MEM_HEAP_FOR_RECV_SYS);
400 #else /* !UNIV_HOTBACKUP */
401 recv_sys->heap = mem_heap_create(256);
402 recv_is_from_backup = TRUE;
403 #endif /* !UNIV_HOTBACKUP */
```

make flush list: see the details later

# Recovery system

- log/log0recv.cc:380, recv\_sys\_init()

```
404
405 /* Set appropriate value of recv_n_pool_free_frames. */
406 if (buf_pool_get_curr_size() >= (10 * 1024 * 1024)) {
407     /* Buffer pool of size greater than 10 MB. */
408     recv_n_pool_free_frames = 512;
409 }
410
411 recv_sys->buf = static_cast<byte*>(ut_malloc(RECV_PARSING_BUF_SIZE));
412 recv_sys->len = 0;
413 recv_sys->recovered_offset = 0;
414
415 recv_sys->addr_hash = hash_create(available_memory / 512);
416 recv_sys->n_addrs = 0;
417
418 recv_sys->apply_log_recs = FALSE;
419 recv_sys->apply_batch_on = FALSE;
420
421 recv_sys->last_block_buf_start = static_cast<byte*>(
422     mem_alloc(2 * OS_FILE_LOG_BLOCK_SIZE));
423
424 recv_sys->last_block = static_cast<byte*>(ut_align(
425     recv_sys->last_block_buf_start, OS_FILE_LOG_BLOCK_SIZE));
426
427 recv_sys->found_corrupt_log = FALSE;
428
429 recv_max_page_lsn = 0;
430
431 /* Call the constructor for recv_sys_t::dblwr member */
432 new (&recv_sys->dblwr) recv_dblwr_t();
433
434 mutex_exit(&(recv_sys->mutex));
435 }
```

make recovery system buffer pool for log scan

Init default variables and make buffers

create double write system

# **OPEN DATA FILES AND LOG FILE**

# check data files

- `srv/srv0start.cc:2059`

```
2059 err = open_or_create_data_files(&create_new_db,  
2060 #ifdef UNIV_LOG_ARCHIVE  
2061     &min_arch_log_no, &max_arch_log_no,  
2062 #endif /* UNIV_LOG_ARCHIVE */  
2063     &min_flushed_lsn, &max_flushed_lsn,  
2064     &sum_of_new_sizes);  
2065 if (err == DB_FAIL) {  
2066  
2067     ib_logf(IB_LOG_LEVEL_ERROR,  
2068         "The system tablespace must be writable!");  
2069  
2070     return(DB_ERROR);  
2071 } else if (err != DB_SUCCESS) {  
2072  
2073     ib_logf(IB_LOG_LEVEL_ERROR,  
2074         "Could not open or create the system tablespace. If "  
2075         "you tried to add new data files to the system "  
2076         "tablespace, and it failed here, you should now "  
2077         "edit innodb_data_file_path in my.cnf back to what "  
2078         "it was, and remove the new ibdata files InnoDB "  
2079         "created in this failed attempt. InnoDB only wrote "  
2080         "those files full of zeros, but did not yet use "  
2081         "them in any way. But be careful: do not remove "  
2082         "old data files which contain your precious data!");  
2083  
2084     return(err);  
2085 }  
2086 }
```

Open data files

open data files error  
handling

# check data files

- `srv/srv0start.cc:759, open_or_create_data_files()`
  - Skip

```
989      /* This is the earliest location where we can load
990      the double write buffer. */
991      if (i == 0) {
992          buf_dblwr_init_or_load_pages(
993              files[i], srv_data_file_names[i], true);
994      }
```

init or load double  
write pages

# check log files

- `srv/srv0start.cc`

```
2119     for (i = 0; i < SRV_N_LOG_FILES_MAX; i++) {
2120         os_offset_t size;
2121         os_file_stat_t stat_info;
2122
2123         sprintf(logfilename + dirnamelen,
2124             "ib_logfile%u", i);
2125
2126         err = os_file_get_status(
2127             logfilename, &stat_info, false);
2128
2129         if (err == DB_NOT_FOUND) {
2130             if (i == 0) {
2131                 if (max_flushed_lsn
2132                     != min_flushed_lsn) {
2133                     ib_logf(IB_LOG_LEVEL_ERROR,
2134                         "Cannot create"
2135                         " log files because"
2136                         " data files are"
2137                         " corrupt or"
2138                         " not in sync"
2139                         " with each other");
2140                     return(DB_ERROR);
2141                 }

```

get log file status

if log file not found, then create  
new one

# check log files

- `srv/srv0start.cc:2189`

```
2185     if (!srv_file_check_mode(logfilename)) {
2186         return(DB_ERROR);
2187     }
2188
2189     err = open_log_file(&files[i], logfilename, &size);
2190
2191     if (err != DB_SUCCESS) {
2192         return(err);
2193     }
2194
2195     ut_a(size != (os_offset_t) -1);
2196
2197     if (size & ((1 << UNIV_PAGE_SIZE_SHIFT) - 1)) {
2198         ib_logf(IB_LOG_LEVEL_ERROR,
2199             "Log file %s size "
2200             UINT64PF " is not a multiple of"
2201             " innodb_page_size",
2202             logfilename, size);
2203         return(DB_ERROR);
2204     }
```

get log file status



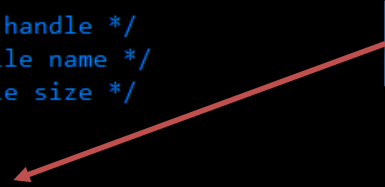


# check log files

- `srv/srv0start.cc:3219, open_log_file()`

```
726 /**
727 Opens a log file.
728 @return DB_SUCCESS or error code */
729 static __attribute__((nonnull, warn_unused_result))
730 dberr_t
731 open_log_file(
732 /*=====*/
733   os_file_t*   file, /*!< out: file handle */
734   const char*  name, /*!< in: log file name */
735   os_offset_t* size) /*!< out: file size */
736 {
737   ibool ret;
738
739   *file = os_file_create(innodb_file_log_key, name,
740                         OS_FILE_OPEN, OS_FILE_AIO,
741                         OS_LOG_FILE, &ret);
742   if (!ret) {
743     ib_logf(IB_LOG_LEVEL_ERROR, "Unable to open '%s'", name);
744     return(DB_ERROR);
745   }
746
747   *size = os_file_get_size(*file);
748
749   ret = os_file_close(*file);
750   ut_a(ret);
751   return(DB_SUCCESS);
752 }
```

create mode = OS\_FILE\_OPEN



# check log files

- `srv/srv0start.cc:2230`

```
2226  /* Create the in-memory file space objects. */
2227
2228  sprintf(logfilename + dirnamelen, "ib_logfile%u", 0);
2229
2230  fil_space_create(logfilename,
2231                  SRV_LOG_SPACE_FIRST_ID,
2232                  fsp_flags_set_page_size(0, UNIV_PAGE_SIZE),
2233                  FIL_LOG);
2234
2235  ut_a(fil_validate());
2236
2237  /* srv_log_file_size is measured in pages; if page size is 16KB,
2238  then we have a limit of 64TB on 32 bit systems */
2239  ut_a(srv_log_file_size <= ULINT_MAX);
2240
2241  for (unsigned j = 0; j < i; j++) {
2242      sprintf(logfilename + dirnamelen, "ib_logfile%u", j);
2243
2244      if (!fil_node_create(logfilename,
2245                          (ulint) srv_log_file_size,
2246                          SRV_LOG_SPACE_FIRST_ID, FALSE)) {
2247          return(DB_ERROR);
2248      }
```

file system create for log files

file node create for log files

# Open log and system table space

- `srv/srv0start.cc:2267,`  
`fil_open_log_and_system_tablespace_files()`

```
2262 files_checked:
2263     /* Open all log files and data files in the system
2264     tablespace: we keep them open until database
2265     shutdown */
2266
2267     fil_open_log_and_system_tablespace_files();
2268
2269     err = srv_undo_tablespaces_init(
2270         create_new_db,
2271         srv_undo_tablespaces,
2272         &srv_undo_tablespaces_open);
2273
2274     /* If the force recovery is set very high then we carry on regardless
2275     of all errors. Basically this is fingers crossed mode. */
2276
2277     if (err != DB_SUCCESS
2278         && srv_force_recovery < SRV_FORCE_NO_UNDO_LOG_SCAN) {
2279
2280         return(err);
2281     }
```

file node create for log files



# Open log and system table space

- fil/fil0fil.cc:1708,  
fil\_open\_log\_and\_system\_tablespace\_files()

```
1706 UNIV_INTERN
1707 void
1708 fil_open_log_and_system_tablespace_files(void)
1709 /*=====*/
1710 {
1711     fil_space_t*   space;
1712
1713     mutex_enter(&fil_system->mutex);
1714
1715     for (space = UT_LIST_GET_FIRST(fil_system->space_list);
1716          space != NULL;
1717          space = UT_LIST_GET_NEXT(space_list, space)) {
1718
1719         fil_node_t* node;
1720
1721         if (fil_space_belongs_in_lru(space)) {
1722             continue;
1723         }
1724     }
```

# Open log and system table space

- `fil/fil0fil.cc:1708,`  
`fil_open_log_and_system_tablespace_files()`

```
1726     for (node = UT_LIST_GET_FIRST(space->chain);
1727          node != NULL;
1728          node = UT_LIST_GET_NEXT(chain, node)) {
1729
1730         if (!node->open) {
1731             if (!fil_node_open_file(node, fil_system,
1732                                     space)) {
1733                 /* This func is called during server's
1734                  * startup. If some file of log or system
1735                  * tablespace is missing, the server
1736                  * can't start successfully. So we should
1737                  * assert for it. */
1738                 ut_a(0);
1739             }
1740         }
1741     }
1742 }
1743
1744 }
1745
1746 }
1747
1748 mutex_exit(&fil_system->mutex);
1749 }
```

file node open



# **INIT UNDO TABLESPACE**

# Undo table space init

- `srv/srv0start.cc:2269,`  
`srv_undo_tablespace_init()`

```
2262 files_checked:
2263     /* Open all log files and data files in the system
2264     tablespace: we keep them open until database
2265     shutdown */
2266
2267     fil_open_log_and_system_tablespace_files();
2268
2269     err = srv_undo_tablespace_init(
2270         create_new_db,
2271         srv_undo_tablespace,
2272         &srv_undo_tablespace_open);
2273
2274     /* If the force recovery is set very high then we carry on regardless
2275     of all errors. Basically this is fingers crossed mode. */
2276
2277     if (err != DB_SUCCESS
2278         && srv_force_recovery < SRV_FORCE_NO_UNDO_LOG_SCAN) {
2279
2280         return(err);
2281     }
```

undo table space init

# Undo table space init

- `srv/srv0start.cc:1272,`  
`srv_undo_tablespace_init()`

```
1270 static
1271 dberr_t
1272 srv_undo_tablespaces_init(
1273 /*=====*/
1274     ibool    create_new_db,    /*!< in: TRUE if new db being
1275                                created */
1276     const uint n_conf_tablespace, /*!< in: configured undo
1277                                tablespaces */
1278     uint*     n_opened) /*!< out: number of UNDO
1279                                tablespaces successfully
1280                                discovered and opened */
1281 {
1282     uint    i;
1283     dberr_t  err = DB_SUCCESS;
1284     uint    prev_space_id = 0;
1285     uint    n_undo_tablespace;
1286     uint    undo_tablespace_ids[TRX_SYS_N_RSEGS + 1];
1287
1288     *n_opened = 0;
1289
1290     ut_a(n_conf_tablespace <= TRX_SYS_N_RSEGS);
1291
1292     memset(undo_tablespace_ids, 0x0, sizeof(undo_tablespace_ids));
```



# Undo table space init

- `srv/srv0start.cc:1272,`  
`srv_undo_tablespace_init()`

```
1294  /* Create the undo spaces only if we are creating a new
1295  instance. We don't allow creating of new undo tablespaces
1296  in an existing instance (yet). This restriction exists because
1297  we check in several places for SYSTEM tablespaces to be less than
1298  the min of user defined tablespace ids. Once we implement saving
1299  the location of the undo tablespaces and their space ids this
1300  restriction will/should be lifted. */
1301
1302  for (i = 0; create_new_db && i < n_conf_tablespaces; ++i) {
1303      char name[OS_FILE_MAX_PATH];
1304
1305      ut_snprintf(
1306          name, sizeof(name),
1307          "%s%cundo%03lu",
1308          srv_undo_dir, SRV_PATH_SEPARATOR, i + 1);
1309
1310      /* Undo space ids start from 1. */
1311      err = srv_undo_tablespace_create(
1312          name, SRV_UNDO_TABLESPACE_SIZE_IN_PAGES);
1313
1314      if (err != DB_SUCCESS) {
1315
1316          ib_logf(IB_LOG_LEVEL_ERROR,
1317              "Could not create undo tablespace '%s'.",
1318              name);
1319
1320          return(err);
1321      }
1322  }
```

create new db case, skip

# Undo table space init

- `srv/srv0start.cc:1272, srv_undo_tablespace_init()`

```
1347 for (i = 0; i < n_undo_tablespaces; ++i) {
1348     char name[OS_FILE_MAX_PATH];
1349
1363     /* Undo space ids start from 1. */
1364
1365     err = srv_undo_tablespace_open(name, undo_tablespace_ids[i]);
1366
1367     if (err != DB_SUCCESS) {
1368
1369         ib_logf(IB_LOG_LEVEL_ERROR,
1370             "Unable to open undo tablespace '%s'.", name);
1371
1372         return(err);
1373     }
1374
1375     prev_space_id = undo_tablespace_ids[i];
1376
1377     ++*n_opened;
1378 }
```

open undo tablespace

# **START TRANSACTION SYSTEM**

# Create transaction system

- `srv/srv0start.cc:2291, trx_sys_create()`

```
2282
2283  /* Initialize objects used by dict stats gathering thread, which
2284     can also be used by recovery if it tries to drop some table */
2285  if (!srv_read_only_mode) {
2286     dict_stats_thread_init();
2287  }
2288
2289  trx_sys_file_format_init();
2290
2291  trx_sys_create();
2292
2293  if (create_new_db) {
2294
2295     ut_a(!srv_read_only_mode);
2296
2297     mtr_start(&mtr);
2298
2299     fsp_header_init(0, sum_of_new_sizes, &mtr);
2300
2301     mtr_commit(&mtr);
2302
```

create transaction system

# Create transaction system

- `trx/trx0sys.cc:587, trx_sys_create()`

```
583 /**
584  Creates the trx_sys instance and initializes ib_bh and mutex. */
585 UNIV_INTERN
586 void
587 trx_sys_create(void)
588 /*=====*/
589 {
590     ut_ad(trx_sys == NULL);
591
592     trx_sys = static_cast<trx_sys_t*>(mem_zalloc(sizeof(*trx_sys)));
593
594     mutex_create(trx_sys_mutex_key, &trx_sys->mutex, SYNC_TRX_SYS);
595 }
596
```

# Do recovery

- `srv/srv0start.cc:2404,`  
`recv_recovery_from_checkpoint_start()`

```
2401  /* We always try to do a recovery, even if the database had
2402  been shut down normally: this is the normal startup path */
2403
2404  err = recv_recovery_from_checkpoint_start(
2405      LOG_CHECKPOINT, LSN_MAX,
2406      min_flushed_lsn, max_flushed_lsn);
2407
2408  if (err != DB_SUCCESS) {
2409
2410      return(DB_ERROR);
2411  }
2412
```

recovery from checkpoint:  
see details later

# Initialize transaction system

- `srv/srv0start.cc:2425,`  
`trx_sys_init_at_db_start()`

see details later

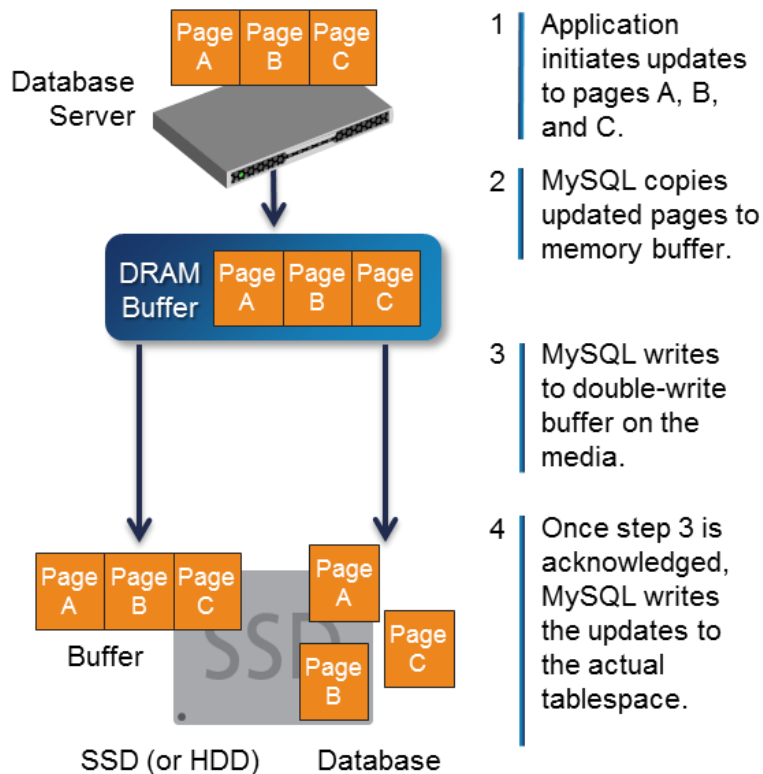
```
2425  ib_bh = trx_sys_init_at_db_start();  
2426  n_recovered_trx = UT_LIST_GET_LEN(trx_sys->rw_trx_list);  
2427  
2428  /* The purge system needs to create the purge view and  
2429     therefore requires that the trx_sys is initied. */  
2430  
2431  trx_purge_sys_create(srv_n_purge_threads, ib_bh);  
2432  
2433  /* recv_recovery_from_checkpoint_finish needs trx lists which  
2434     are initialized in trx_sys_init_at_db_start(). */  
2435  
2436  recv_recovery_from_checkpoint_finish();  
2437
```

# **CREATE DOUBLE WRITE BUFFER**

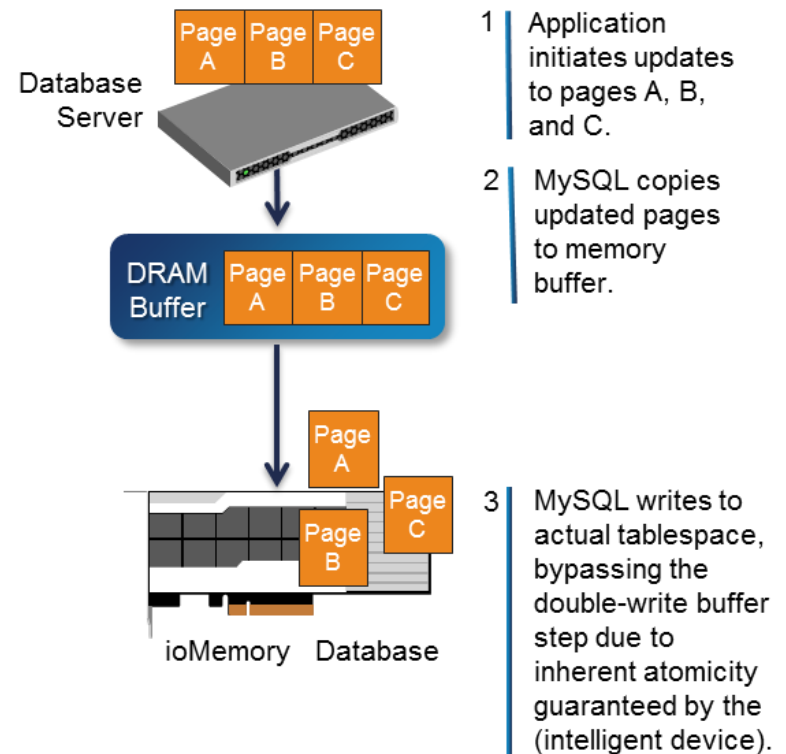


# Double write buffer

## Traditional MySQL Writes



## MySQL with Atomic Writes



[http://www.fusionio.com/load/-media-/29ek9p/imagesLibrary/03\\_atomic\\_writes.png](http://www.fusionio.com/load/-media-/29ek9p/imagesLibrary/03_atomic_writes.png)

# double write create

- `srv/srv0start.cc:2605,`  
`buf_dblwr_create()`

```
2601
2602  if (buf_dblwr == NULL) {
2603      /* Create the doublewrite buffer to a new tablespace */
2604
2605      buf_dblwr_create();
2606  }
2607
2608  /* Here the double write buffer has already been created and so
2609  any new rollback segments will be allocated after the double
2610  write buffer. The default segment should already exist.
2611  We create the new segments only if it's a new database or
2612  the database was shutdown cleanly. */
2613
2614  /* Note: When creating the extra rollback segments during an upgrade
2615  we violate the latching order, even if the change buffer is empty.
2616  We make an exception in sync0sync.cc and check srv_is_being_started
2617  for that violation. It cannot create a deadlock because we are still
2618  running in single threaded mode essentially. Only the IO threads
2619  should be running at this stage. */
```

# double write create

- buf/buf0dblwr.cc:178, buf\_dblwr\_create()

```
173 /**
174  Creates the doublewrite buffer to a new InnoDB installation. The header of the
175  doublewrite buffer is placed on the trx system header page. */
176 UNIV_INTERN
177 void
178 buf_dblwr_create(void)
179 /*=====*/
180 {
181     buf_block_t*   block2;
182     buf_block_t*   new_block;
183     byte*           doublewrite;
184     byte*           fseg_header;
185     ulint page_no;
186     ulint prev_page_no;
187     ulint i;
188     mtr_t mtr;
189
190     if (buf_dblwr) {
191         /* Already init'd */
192
193         return;
194     }
195
196 start_again:
197     mtr_start(&mtr);
198     buf_dblwr_being_created = TRUE;
199
200     doublewrite = buf_dblwr_get(&mtr);
```

# double write create

- buf/buf0dblwr.cc:178, buf\_dblwr\_create()

```
173 /**
174  Creates the doublewrite buffer to a new InnoDB installation. The header of the
175  doublewrite buffer is placed on the trx system header page. */
176 UNIV_INTERN
177 void
178 buf_dblwr_create(void)
179 /*=====*/
180 {
181     buf_block_t*   block2;
182     buf_block_t*   new_block;
183     byte*          doublewrite;
184     byte*          fseg_header;
185     ulint page_no;
186     ulint prev_page_no;
187     ulint i;
188     mtr_t mtr;
189
190     if (buf_dblwr) {
191         /* Already init'd */
192
193         return;
194     }
195
196     start_again:
197     mtr_start(&mtr);
198     buf_dblwr_being_created = TRUE;
199
200     doublewrite = buf_dblwr_get(&mtr);
```

get double write buffer



# double write create

- buf/buf0dblwr.cc:178, buf\_dblwr\_create()

```
202  if (mach_read_from_4(doublewrite + TRX_SYS_DOUBLEWRITE_MAGIC)
203      == TRX_SYS_DOUBLEWRITE_MAGIC_N) {
204      /* The doublewrite buffer has already been created:
205      just read in some numbers */
206
207      buf_dblwr_init(doublewrite);
208
209      mtr_commit(&mtr);
210      buf_dblwr_being_created = FALSE;
211      return;
212  }
213
214  ib_logf(IB_LOG_LEVEL_INFO,
215          "Doublewrite buffer not found: creating new");
216
217  if (buf_pool_get_curr_size()
218      < ((2 * TRX_SYS_DOUBLEWRITE_BLOCK_SIZE
219          + FSP_EXTENT_SIZE / 2 + 100)
220          * UNIV_PAGE_SIZE)) {
221
222      ib_logf(IB_LOG_LEVEL_ERROR,
223              "Cannot create doublewrite buffer: you must "
224              "increase your buffer pool size. Cannot continue "
225              "operation.");
226
227      exit(EXIT_FAILURE);
228  }
```

doublewrite buffer  
already created

start of doublewrite buffer creation:  
skip

# Note

- `mach_read_from_4`
  - read 4 bytes
  - considering architecture specific endian.

```
175 /*****
176 The following function is used to fetch data from 4 consecutive
177 bytes. The most significant byte is at the lowest address.
178 @return uint integer */
179 UNIV_INLINE
180 uint
181 mach_read_from_4(
182 /*=====*/
183  const byte* b) /*!< in: pointer to four bytes */
184 {
185     ut_ad(b);
186     return( ((uint)(b[0]) << 24)
187         | ((uint)(b[1]) << 16)
188         | ((uint)(b[2]) << 8)
189         | (uint)(b[3])
190     );
191 }
```

# double write create

- buf/buf0dblwr.cc:126, buf\_dblwr\_init()

```
122 /**
123  * Creates or initializes the doublewrite buffer at a database start.
124  */
125 static
126 void
127 buf_dblwr_init(
128     byte* doublewrite) /*!< in: pointer to the doublewrite buffer
129                          header on trx sys page */
130 {
131     ulint buf_size;
132
133     buf_dblwr = static_cast<buf_dblwr_t*>(
134         mem_zalloc(sizeof(buf_dblwr_t)));
135
136     /* There are two blocks of same size in the doublewrite
137      * buffer. */
138     buf_size = 2 * TRX_SYS_DOUBLEWRITE_BLOCK_SIZE;
139
140     /* There must be atleast one buffer for single page writes
141      * and one buffer for batch writes. */
142     ut_a(srv_doublewrite_batch_size > 0
143         && srv_doublewrite_batch_size < buf_size);
144
145     mutex_create(buf_dblwr_mutex_key,
146                 &buf_dblwr->mutex, SYNC_DOUBLEWRITE);
147 }
```

size of doublewrite buffer  
= 2 \* EXTENTS (1MB)

# double write create

- buf/buf0dblwr.cc:126, buf\_dblwr\_init()

```
148 buf_dblwr->b_event = os_event_create();
149 buf_dblwr->s_event = os_event_create();
150 buf_dblwr->first_free = 0;
151 buf_dblwr->s_reserved = 0;
152 buf_dblwr->b_reserved = 0;
153
154 buf_dblwr->block1 = mach_read_from_4(
155     doublewrite + TRX_SYS_DOUBLEWRITE_BLOCK1);
156 buf_dblwr->block2 = mach_read_from_4(
157     doublewrite + TRX_SYS_DOUBLEWRITE_BLOCK2);
158
159 buf_dblwr->in_use = static_cast<bool*>(
160     mem_zalloc(buf_size * sizeof(bool)));
161
162 buf_dblwr->write_buf_unaligned = static_cast<byte*>(
163     ut_malloc((1 + buf_size) * UNIV_PAGE_SIZE));
164
165 buf_dblwr->write_buf = static_cast<byte*>(
166     ut_align(buf_dblwr->write_buf_unaligned,
167     UNIV_PAGE_SIZE));
168
169 buf_dblwr->buf_block_arr = static_cast<buf_page_t**>(
170     mem_zalloc(buf_size * sizeof(void*)));
171 }
```

set 2 blocks of double write header

create buffer frame for doublewrite  
(and make it page sized align)

Prepare buf block array



# Q&A

- Any Questions ?

# Reference

- Source Code : MySQL Community Server 5.6.21
- Transaction Model and Locking
  - <https://dev.mysql.com/doc/refman/5.6/en/innoDB-transaction-model.html>
  - <http://dev.cs.uni-magdeburg.de/db/mysql/InnoDB-transaction-model.html>